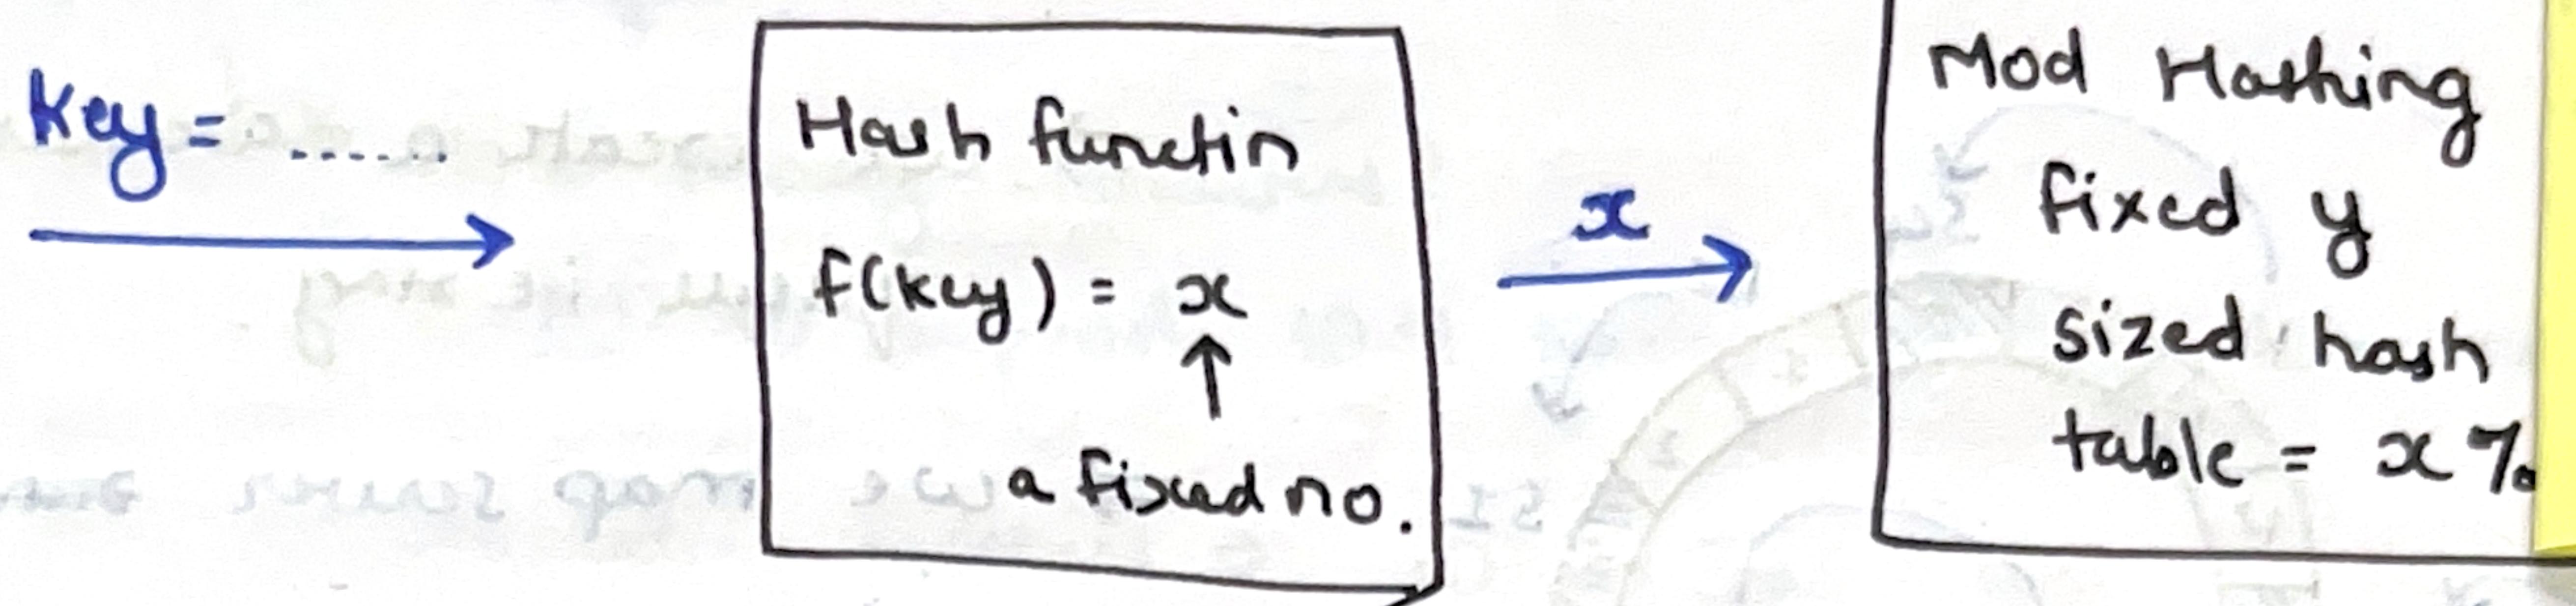


Consistent Hashing



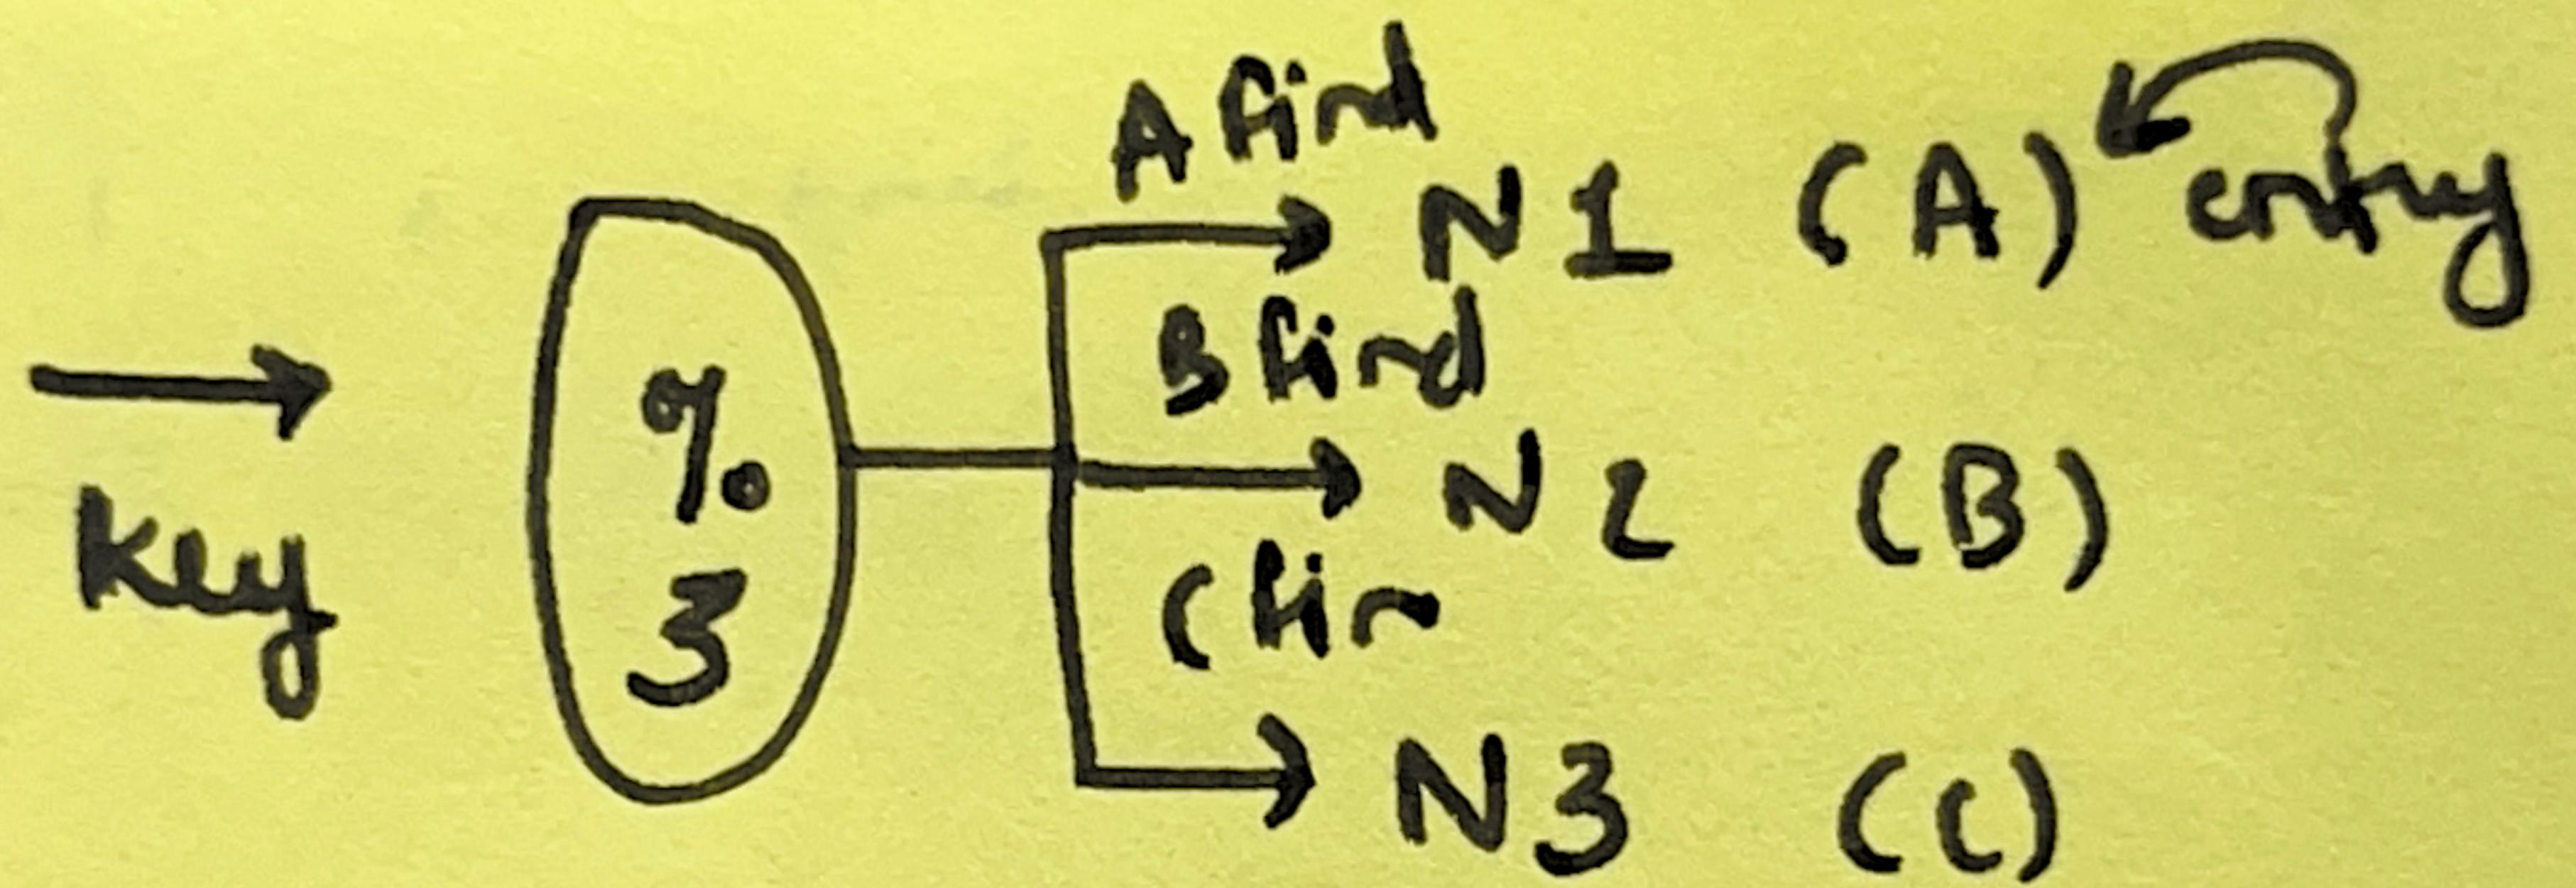
takes key & generates
a fixed no.

does modular
op. on fixed no.

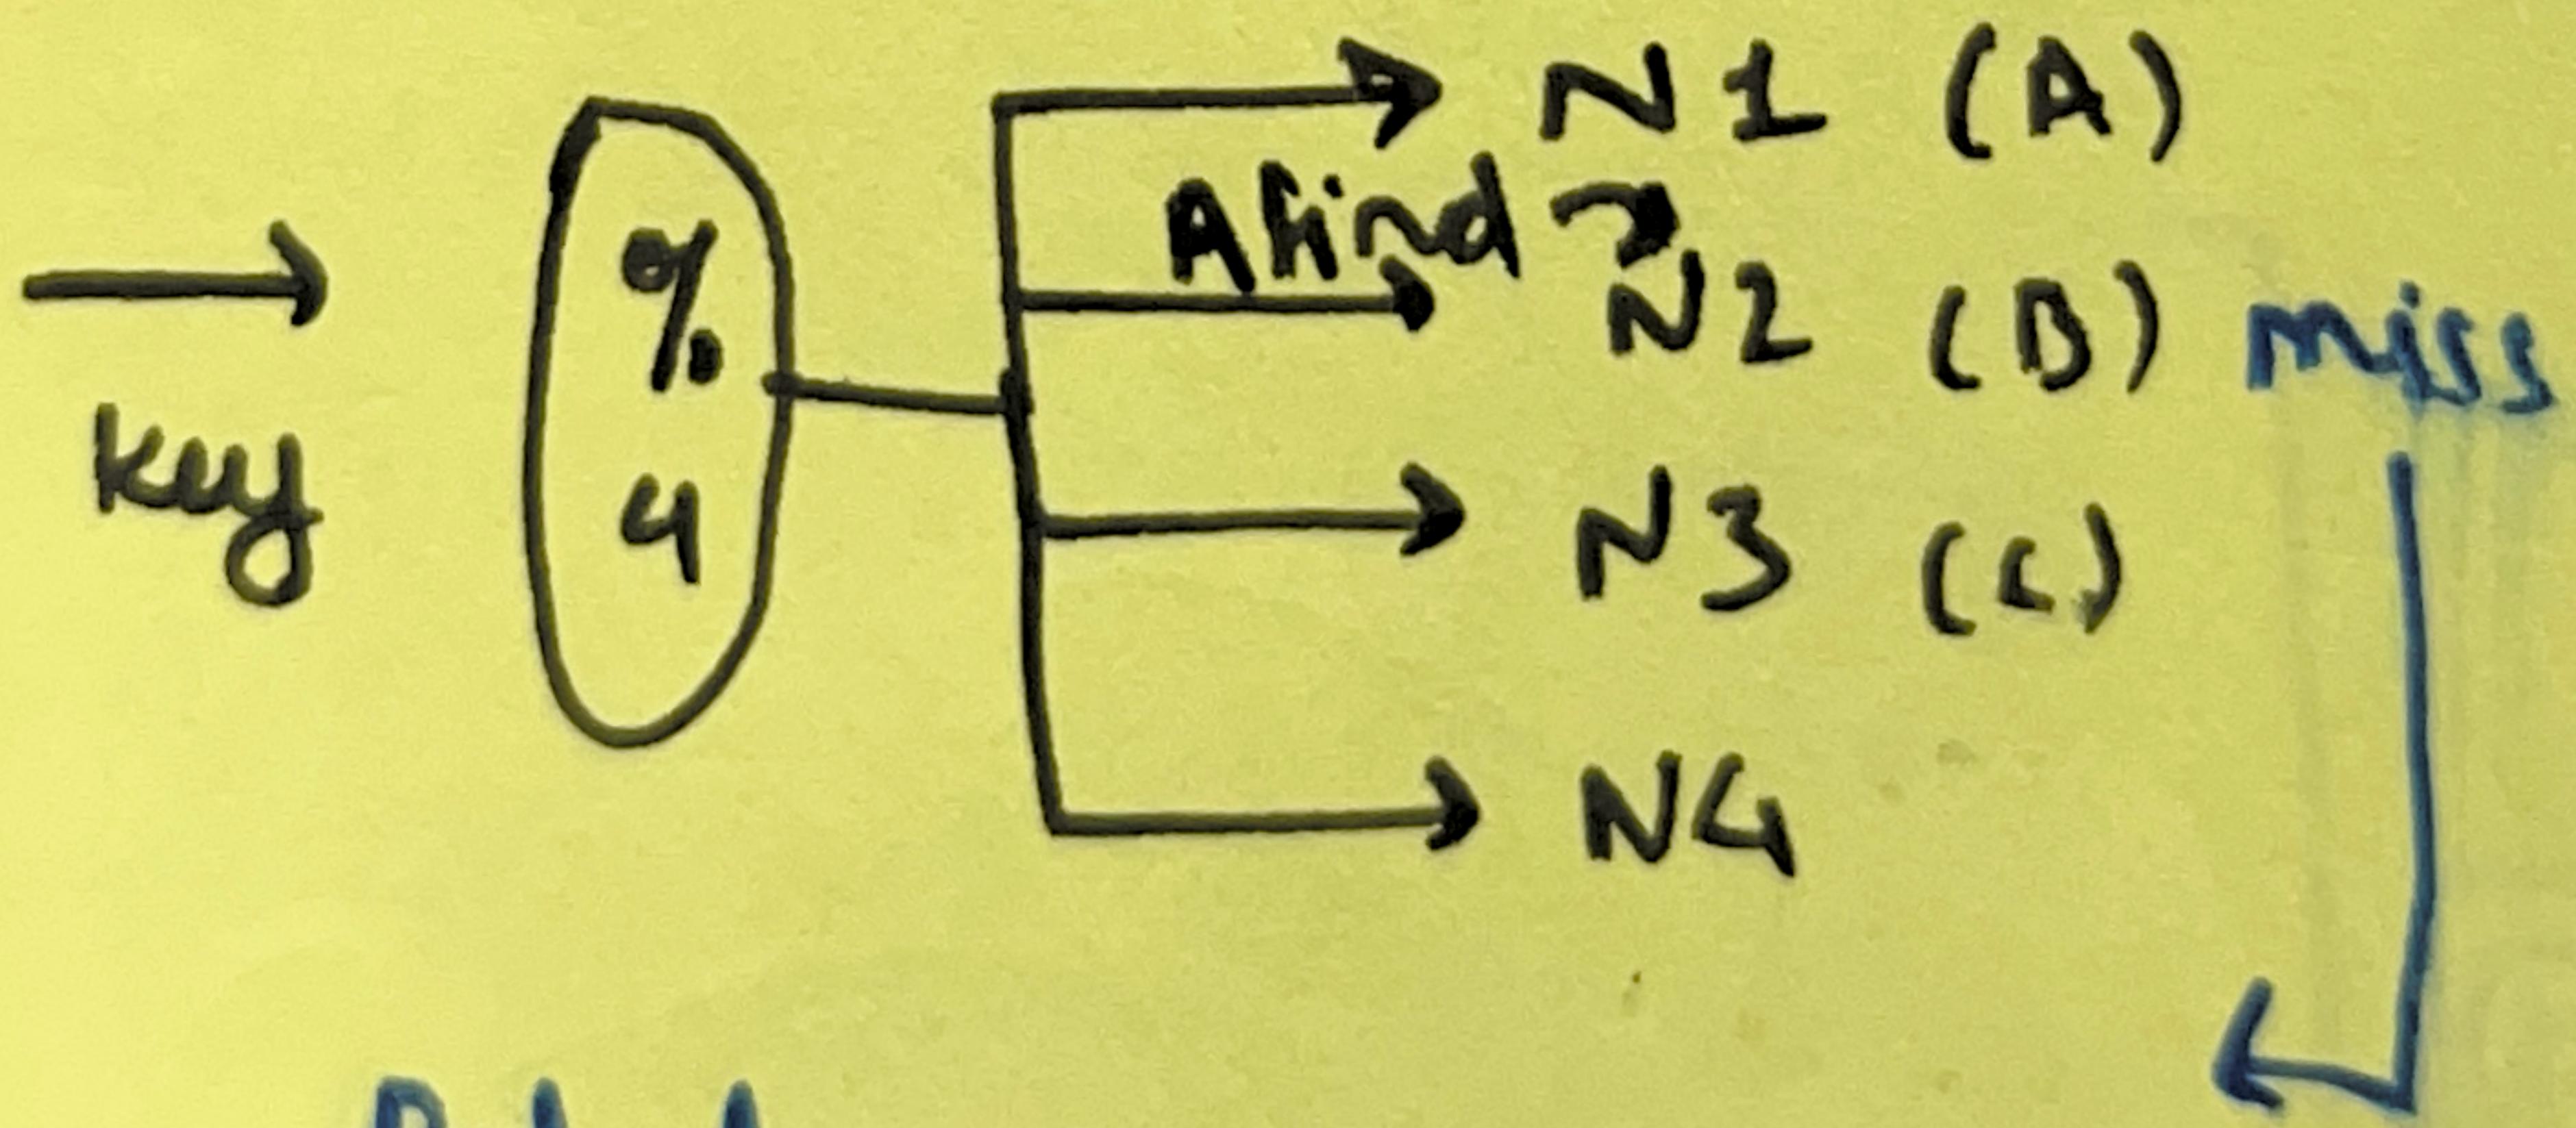
above hashing works perfectly fine when we have fixed size of hash table, But in real world no. of nodes i.e. output of hashing would never be constant.

For Ex: Load Balancing } nodes are never constant, so
DB Sharding } if we keep hash table of same length then newly spawned up servers will always be idle.

* When new DB server is added our hashing function would have mod(y+1): this leads to data being searched in wrong node, thus we would need costly rebalancing if new nodes added.



↓ adding Node



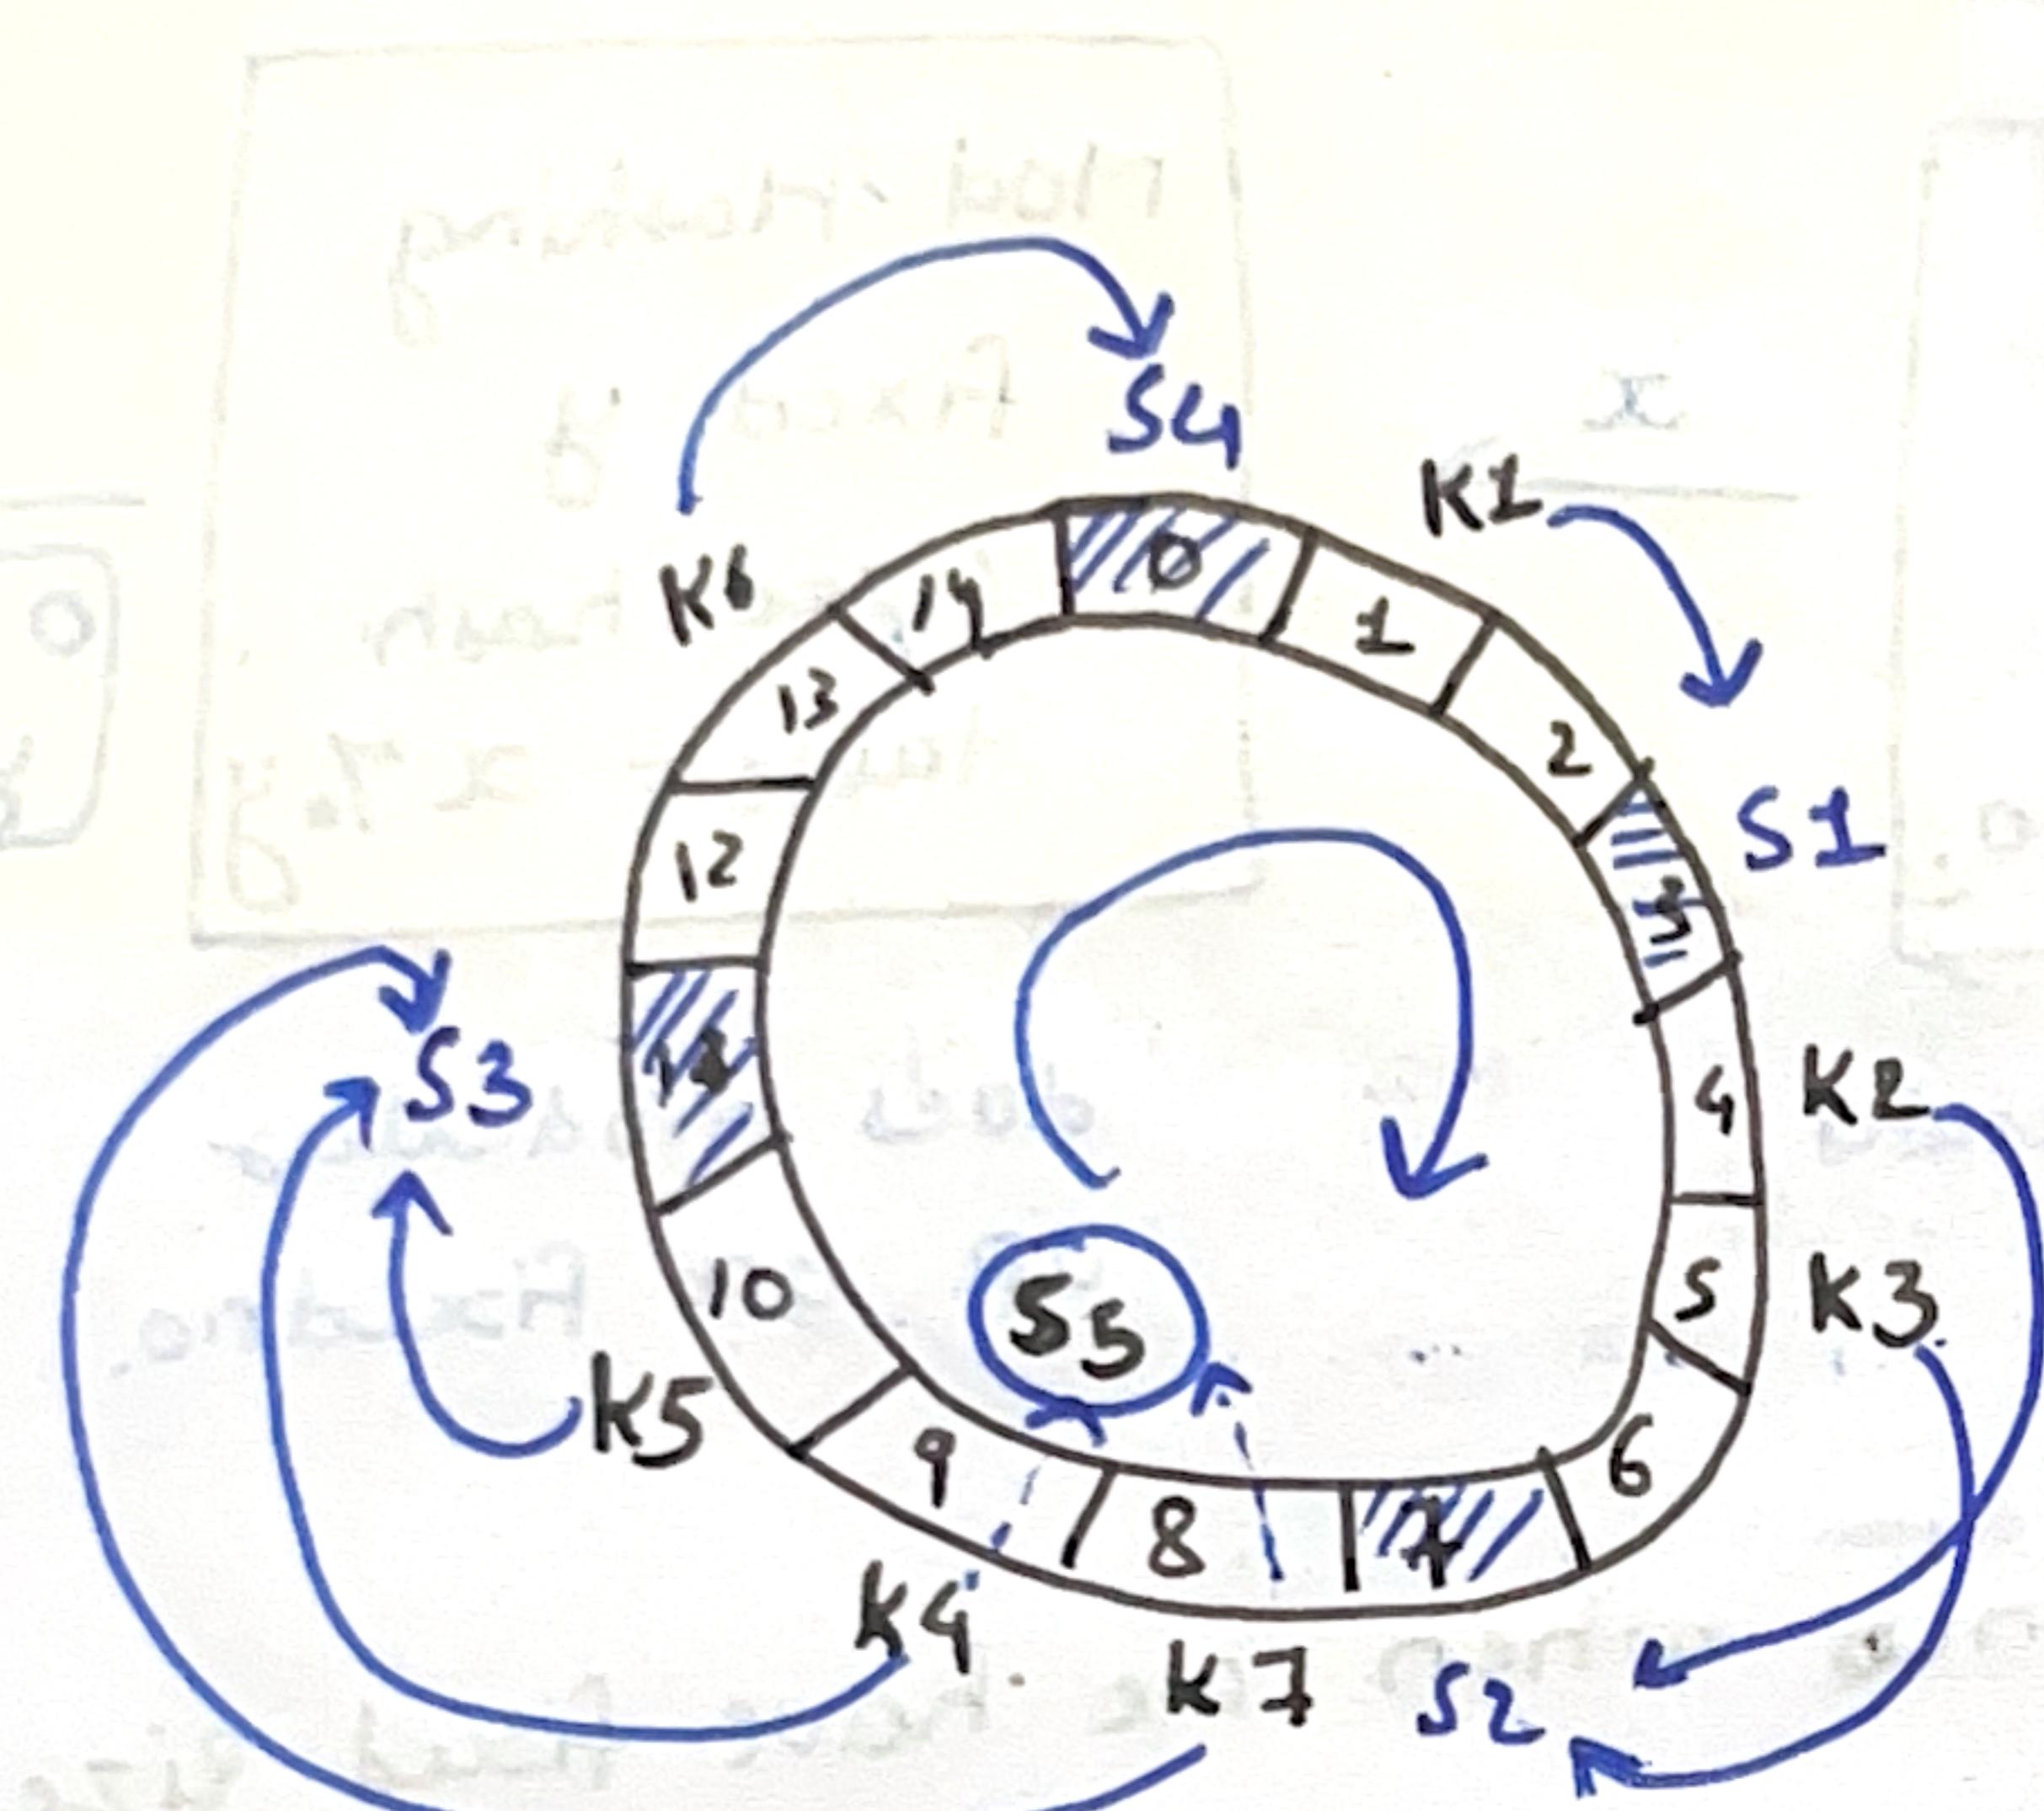
Rebalancing seq.

what does consistency Hashing tend to solve?

⇒ in case of adding/removal of a node(server)

as low as $\frac{1}{n}\%$ of rebalancing is required
High and not more than that.
total nodes

i.e. Reducing Rebalancing among servers.



K_i = Requests

S_i = Server

⇒ we create a circular queue i.e. ring.

⇒ we map server randomly using mod hashing $f(S_i) \% 15$
Hashfn ↑
ring size ↑

⇒ incoming request are also mapped to ring using same mod hashing

⇒ to fulfill the request, we choose nearest server in clockwise direction.

⇒ adding new server can lead to rebalancing of very few req. [added through same mod hashing]

Tip: use replicas of same servers over the ring to avoid request getting congested in one server.

Virtual Objects: Avoid disadvantage due to clouting of servers nearby due to Hashing.